Evan Lee

(925) 791-9632 | exl2813@berkeley.edu | linkedin.com/in/evxlee | github.com/EvxLee | evxlee.github.io

Education

University of California, Berkeley — Berkeley, CA

Bachelor's in Data Science, Bachelor's in Economics, Minor in Computer Science — *Expected Graduation: May 2029*

Relevant Coursework: Principles and Techniques of Data Science, Data Structures, Business Analytics, Discrete Mathematics and Probability Theory, Linear Algebra

Experience

Git Machine — Jun 2026 – Aug 2026

Forward Deployed Engineer Intern — *San Francisco, CA*

Git Machine builds software that recovers wasted GPU compute so teams run more workloads on hardware they already own.

- Drove customer discovery across 20 calls with engineers at AI, robotics, and biotech companies, surfacing pain points around wasted GPU compute; qualified 10 into product evaluations and set up 3 founder calls
- A/B tested two value propositions, throughput gains vs cost savings, across 120 prospects at 50+ companies, lifting reply rate from 5% to 28%
- Collaborated with the engineering team to prioritize the top product requirement from discovery calls: a no-install diagnostic that shows prospects how much GPU compute they waste before committing

Ithildin (Berkeley SkyDeck) — Jun 2026 – Aug 2026

AI/ML Engineer Intern — *Berkeley, CA*

Deposition platform that transcribes testimony and flags inconsistencies for litigators and court reporting firms.

- Defined the benchmark for selecting transcription models, prioritizing speaker accuracy over word accuracy to minimize incorrect speaker identification; reduced DER from 20.0% to 12.9%
- Designed the two-pass transcription engine the product is built on, a streaming transcript for quick reference and a full-recording pass for the certified record, after learning litigators review at breaks rather than live

UCSF Yang Lab — May 2026 – Sep 2026

NLP Researcher - advised under Dr. Hanxue Gu — *San Francisco, CA*

- Co-authored a NeurIPS 2026 workshop submission on CohortFlow, an AI agent that finds trial-eligible patients, designed around how clinicians work: criteria change mid-run, cutting correction latency 64% at matched accuracy
- Developed an NLP pipeline extracting 6 clinical-trial metadata fields from biomedical articles; automated prompt optimization using DSPy to raise accuracy by 9% on 60 held-out papers

LG NOVA — Feb 2026 – May 2026

Machine Learning Engineer (Contract) — *Santa Clara, CA*

LG Electronics' startup incubator; Atlas is its internal platform tracking every startup, project, and contract.

- Built Atlas's RAG chatbot from zero, letting analysts ask portfolio questions in natural language and get accurate, relevant answers instead of manually searching thousands of records across multiple databases
- Automated the quarterly cycle report with AI summaries that classify each startup by industry and vertical, so internal analysts view trends across program cohorts without reading every startup
- Benchmarked the chatbot query quality with Precision@5 on 100+ analyst questions to ensure the internal team could trust its answers

FactGrid Cuneiform — Jan 2026 – Apr 2026

NLP Researcher - advised under Dr. Adam Anderson — *Berkeley, CA*

- Fine-tuned a ByT5 model to correct OCR errors in scanned Akkadian transliterations, cutting character error rate by 75% so historians retrieve the right words when querying FactGrid's open cuneiform database

Projects

Data Matchmaker Benchmark AI Agent | *AgentX – AgentBeats Competition (Berkeley RDI)* Dec 2025 – Jan 2026

- Achieved 2nd place in the OpenAI-sponsored Finance Agent track; deployed an evaluator agent that scores other AI agents' accuracy on data integration tasks

Equity-Focused College Matcher (EquiMatch) — Nov 2025

- Placed 3rd of 50+ teams at UC Berkeley's 7th Annual Datathon for Social Good; built a college recommender app for underserved students, ranking 6,000+ institutions on equity metrics

Skills

Product: Customer Discovery, A/B Testing, Go-to-Market, Product Requirements, Prioritization, Evaluation Benchmarks
Technical: Python, SQL, R, Java, PyTorch, scikit-learn, Hugging Face, RAG, Tableau, PostgreSQL, Docker, AWS, Git